

The University of Sheffield

ACS61012 Machine Vision



Leander Stephen Desouza

Registration Number: 230118120

Department of Automatic Control Systems
Engineering

Contents

| | |
|---|-----------|
| 1 Task 1 - Introduction to Machine Vision | 1 |
| 1.1 Part I: Introduction to Image Processing | 1 |
| 1.1.1 RGB Histogram | 1 |
| 1.2 Part II: Edge Detection and Segmentation of Static Objects | 3 |
| 1.2.1 Static Objects Segmentation by Edge Detection | 3 |
| 2 Task 2 - Object Motion Detection and Tracking | 5 |
| 2.1 Corner analysis using Optical Flow | 5 |
| 2.2 Optical Flow test for GingerBreadMan | 8 |
| 3 Task 3 - Automatic Detection of Moving Objects in a Sequence of Video Frames | 10 |
| 3.1 Frame Differencing Approach | 10 |
| 3.2 Gaussian Mixture Models | 11 |
| 4 Task 4 - Robot Treasure Hunting | 13 |
| 4.1 Easy Case | 13 |
| 4.2 Medium Case | 15 |
| 4.3 Hard Case | 16 |
| 5 Task 5 - Convolutional Neural Networks for Image Classification | 18 |
| 5.1 Accuracy of Default LeNet5 | 18 |
| 5.2 Precision, Recall and F1 Score of Default LeNet5 Model | 19 |
| 5.3 Improvement of the LeNet5 Model | 20 |
| 5.4 Ethics in Computer Vision | 23 |
| 5.4.1 Positives of using Computer Vision | 23 |
| 5.4.2 Challenges and Mitigation of using Computer Vision | 24 |
| 5.4.3 Equality and Diversity in Computer Vision | 24 |
| 6 Appendix | 25 |

List of Figures

| | | |
|----|--|----|
| 1 | RGB Histogram of Mallard indicating low resolution and high intensity | 2 |
| 2 | RGB histogram of a dog indicating high resolution and low intensity | 2 |
| 3 | Comparison of various edge detection algorithms | 4 |
| 4 | Corners detected in the red square sample image | 5 |
| 5 | Estimated Trajectory using Lucas Kanade vs Ground Truth Data | 6 |
| 6 | RMSE comparison throughout the video frames | 7 |
| 7 | Corner points in the GingerBreadMan | 8 |
| 8 | Flow vectors superimposed on the second image of GingerBreadMan | 9 |
| 9 | Plot of the flow vectors observed after Lucas-Kanade optical estimation | 9 |
| 10 | Frame Difference with a binary thresholding of 10 and 25, at frame 140 | 11 |
| 11 | GMM with $n_frames = [1, 10]$, and $n_gaussians = [2, 20]$ at frame 140 | 12 |
| 12 | Intermediate steps in the Easy Case | 13 |
| 13 | Intermediate steps in the Medium Case | 15 |
| 14 | Intermediate treasure in Hard Case | 16 |
| 15 | Intermediate steps in Hard Case | 17 |
| 16 | Training Progress on LeNet5 | 19 |
| 17 | Confusion Matrix of the LeNet5 model | 19 |
| 18 | Training Progress on Custom model | 22 |
| 19 | Confusion Matrix of the Custom model | 22 |

1 Task 1 - Introduction to Machine Vision

1.1 Part I: Introduction to Image Processing

1.1.1 RGB Histogram

An elementary way to analyze the image composition is to describe its RGB contribution. This helps us to understand the overall wavelengths in the image and can help in masking out certain features of the image [1]. The basic idea is if a certain area of the image needs to be segmented, its RGB or HSV ranges can be studied, and a suitable mask can be applied.

Analysis

In the following figures, we have compared the RGB histograms of a dog and mallard. For more information on the RGB histogram embedding, refer to Task -1, part -1 of the Appendix. The clear inference is the intensity of the supplied images. That is observed in the shift in the intensity of the graph; the figure 3 has low intensity and hence has the plot shifted towards the origin, whereas the 1 has greater intensity and is away from the origin.

The second inference is in the pixel density of the images. The mallard image has a significantly lower resolution than that of a dog, and hence, the frequency of the pixels is much lower as well.

In addition, there is no clear domination of a particular channel in either image. However, the green colour dominates in figure 1; this is due to the increased pixels in the mallard's head and the green grass. In figure 3, the blue colour dominates as the image seems to be taken at night with a blueish-green tint.

Conclusion

By comparing the plots of the mallard and dog, we understand the differences in the intensity, RGB composition and image resolution. The mallard has a high intensity and lower image resolution and is green dominant, whereas the dog has a lower intensity and higher image resolution and is blue dominant as inferred from the plot.

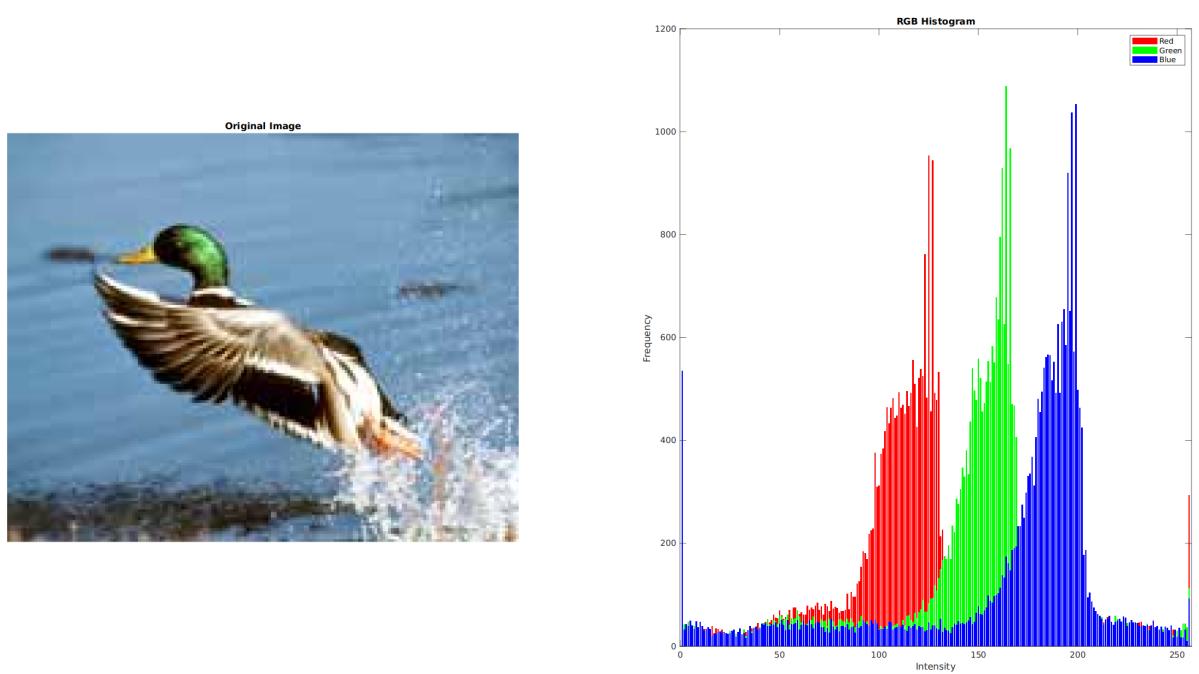


Figure 1: RGB Histogram of Mallard indicating low resolution and high intensity

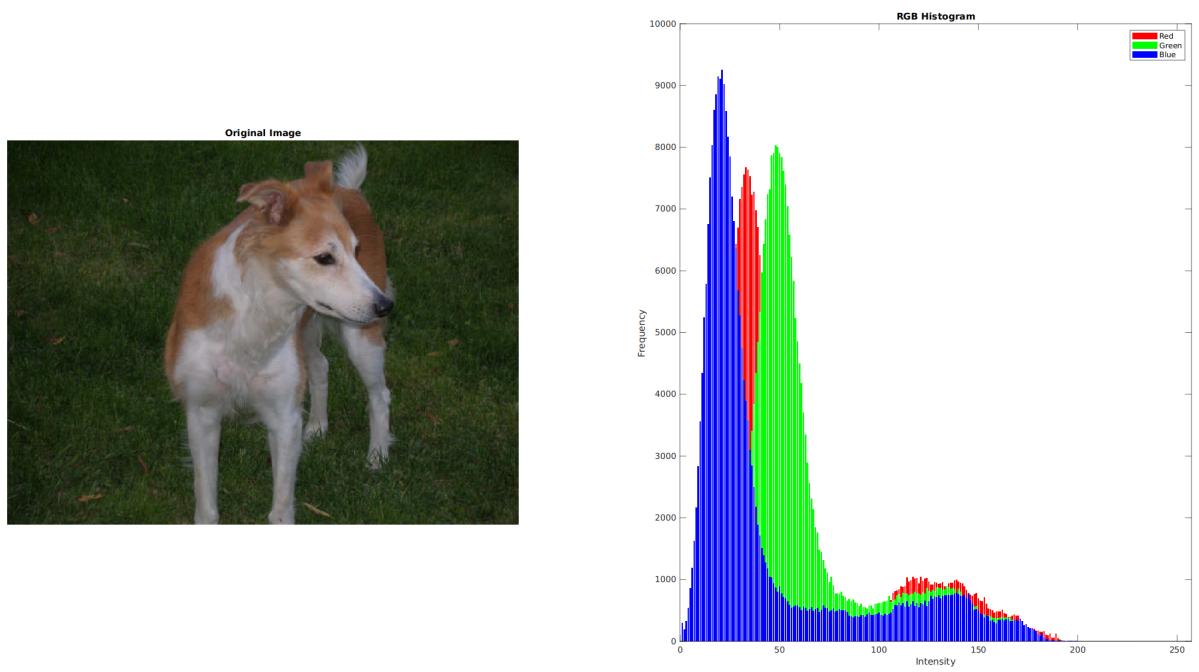


Figure 2: RGB histogram of a dog indicating high resolution and low intensity

1.2 Part II: Edge Detection and Segmentation of Static Objects

1.2.1 Static Objects Segmentation by Edge Detection

Here, we compare three popular edge detection algorithms, namely, Sobel, Canny and Prewitt methods. For all of these detection methods, the input needs to be converted into grayscale images as additional channels provide no data relevant to edge analysis [2]. For additional information about the subplots of the operators, refer to Task -1, part 2 of the Appendix. **Analysis**

1. Sobel Operator:

This method of edge detection focuses on taking gradients of the image intensity function. The derivatives with respect to horizontal and vertical directions are taken and combined to find the edges in each direction. Finally, this is convoluted with the original image to highlight areas where the partial derivative is constant.

This is incorporated in the form of edge detection kernels for each direction:

$$G_x = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$$

$$G_y = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix}$$

2. Canny Operator:

This method relies on finding the maximum intensity gradient of the image and suppressing the unwanted pixels after a full initial scan. Furthermore, a hysteresis thresholding step is applied to differentiate between false positives. Here, the initial step of Gaussian filtering proves to be essential as this filter tends to be more sensitive to noises that are Gaussian in nature.

3. Prewitt Operator:

Similar to the Sobel operator, the Prewitt operator focuses on the approximation of the gradient of the intensity function of the image. However, the key difference lies in the computation of the gradient. This operator is applied on smaller filters that

span across horizontal and vertical directions.

The following are the corresponding kernels for each direction:

$$G_x = \begin{bmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{bmatrix}$$

$$G_y = \begin{bmatrix} -1 & -1 & -1 \\ 0 & 0 & 0 \\ 1 & 1 & 1 \end{bmatrix}$$



Figure 3: Comparison of various edge detection algorithms

Conclusion

Various edge detection methods have been applied to the image of the mallard, as shown. All of the algorithms perform comparatively well. However, even after extensive tuning, there seems to be many false positives in the Canny edge detection image. This may be due to the initial Gaussian filtering step that is involved in the Canny operator, as the filter can have a tendency to nullify some image data as well.

2 Task 2 - Object Motion Detection and Tracking

2.1 Corner analysis using Optical Flow

In this task, we track a specified object of interest throughout its duration in a video frame. We find the optical flow of the vectors of importance and understand the overall trajectory that the object takes in its lifetime [3].

First, we estimate the corner points of the red square provided to us in the lab brief. The procedure used to detect corners here is the Harris Corner Detection Method. The final output is seen in 4, where the corners are detected by the + sign.

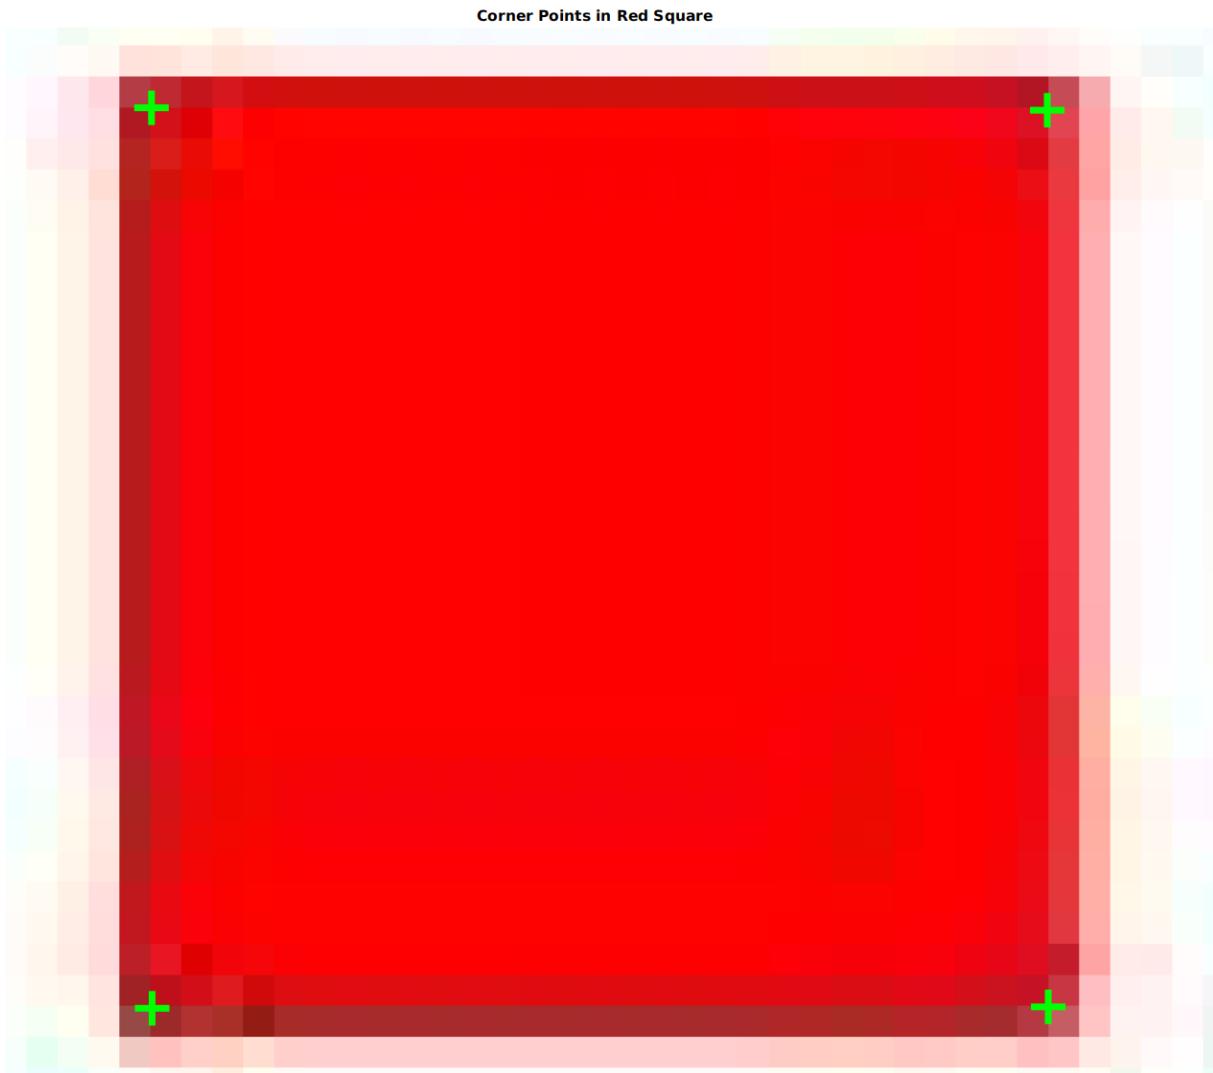


Figure 4: Corners detected in the red square sample image

We then proceed to track the top left corner of the square throughout the video by using the Harris method. The optical flow of the video is estimated using the Lucas

Kanade method, and the estimated new corner is interpolated using the flow calculated. This is then compared with the ground truth data provided while maintaining the RMSE score.

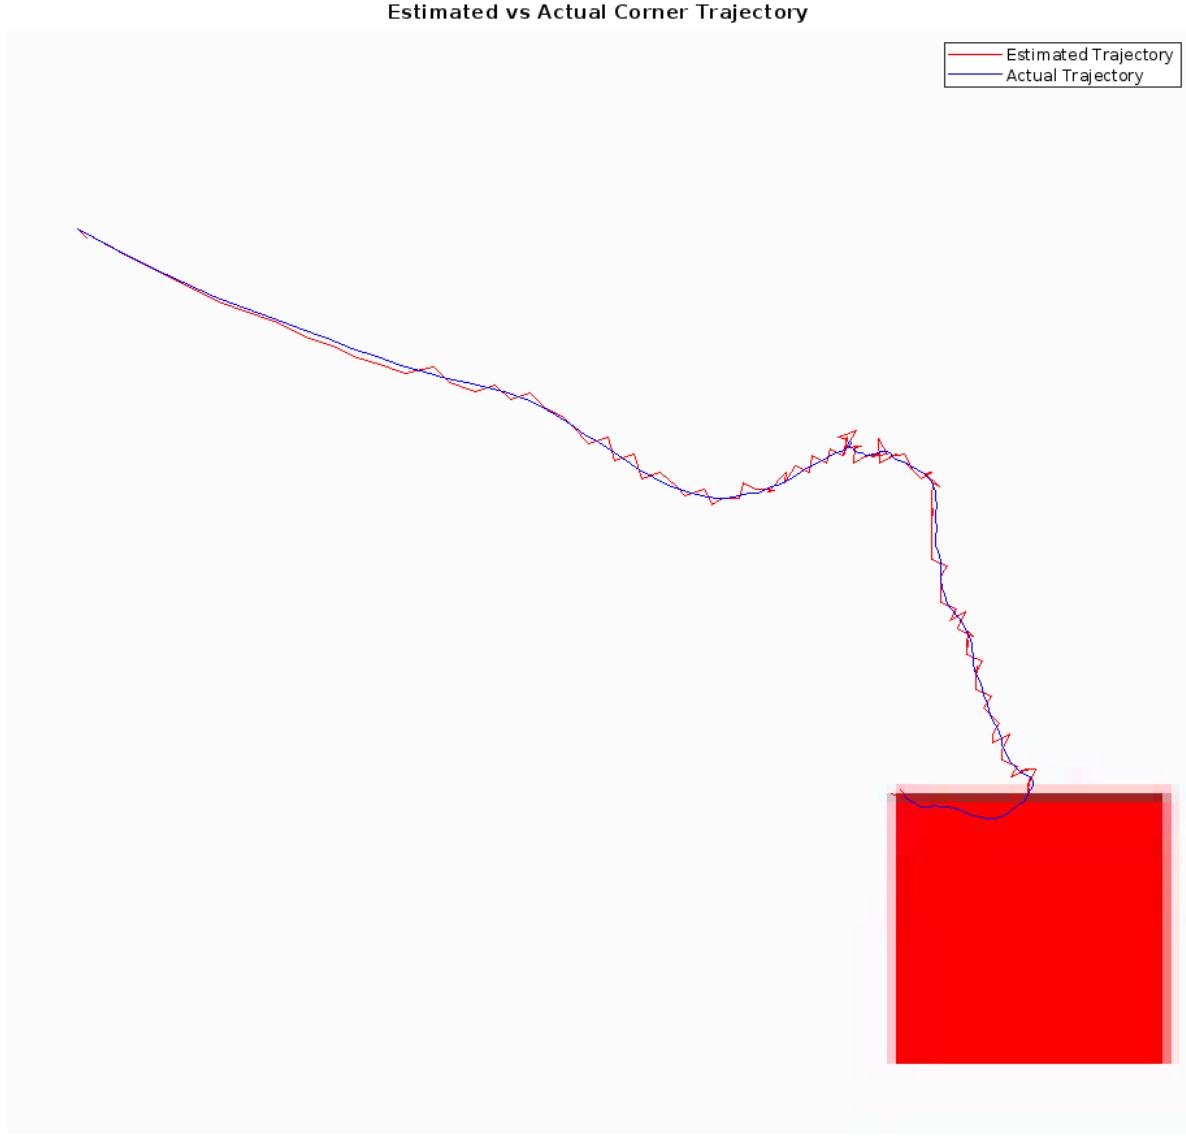


Figure 5: Estimated Trajectory using Lucas Kanade vs Ground Truth Data

As seen in the figure 5, the estimated trajectory closely follows the ground truth with some occasional spikes. This deviation is measured and calculated as follows:

$$RMSE_x = \frac{1}{n-1} \sum_{i=1}^{n-1} \sqrt{(x_{actual}(i+1) - x_{predicted}(i))^2}$$

$$RMSE_y = \frac{1}{n-1} \sum_{i=1}^{n-1} \sqrt{(y_{actual}(i+1) - y_{predicted}(i))^2}$$

$$RMSE_{xy} = \frac{1}{n-1} \sum_{i=1}^{n-1} \sqrt{(x_{actual}(i+1) - x_{predicted}(i))^2 + (y_{actual}(i+1) - y_{predicted}(i))^2}$$

Here, n is the number of frames, and i is each frame in the corresponding video stream. As you can see, the predicted frame is compared with the next future frame in order to determine the stability of the optical flow, hence the figure 6 starts from the second frame. All these formulas have been accordingly in Task 2, part 1 of the Appendix.

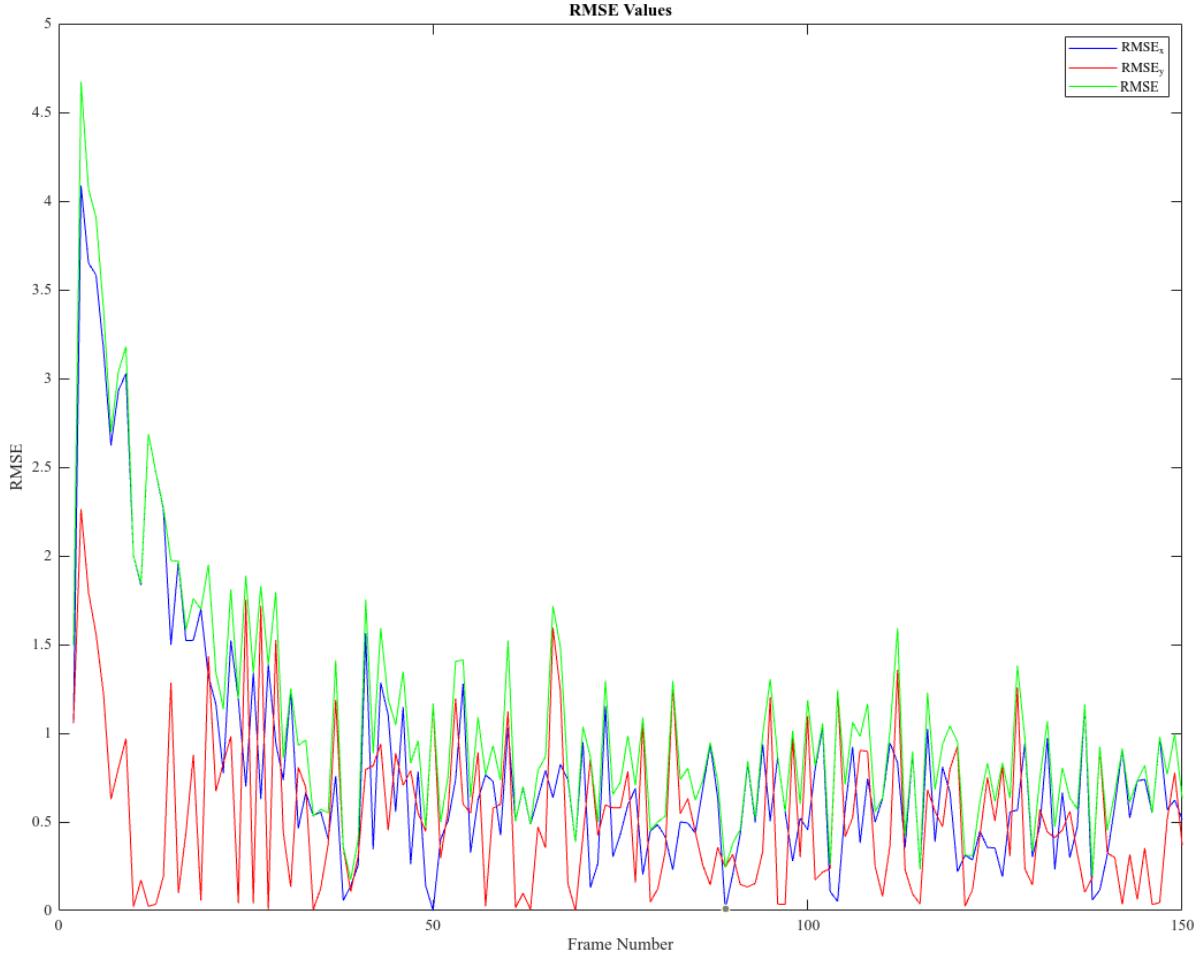


Figure 6: RMSE comparison throughout the video frames

From the figure, the RMSE starts at a high value but effectively converges as the video frames increase, resulting in a better prediction. The overall RMSE values are as follows:

$$RMSE_x = 0.8329$$

$$RMSE_y = 0.5483$$

$$RMSE_{xy} = 1.0891$$

2.2 Optical Flow test for GingerBreadMan

Similarly to the red square, first, we need to track the interesting points in the gingerbread man. Therefore, we extract the corner points using Harris method as seen in figure 7.

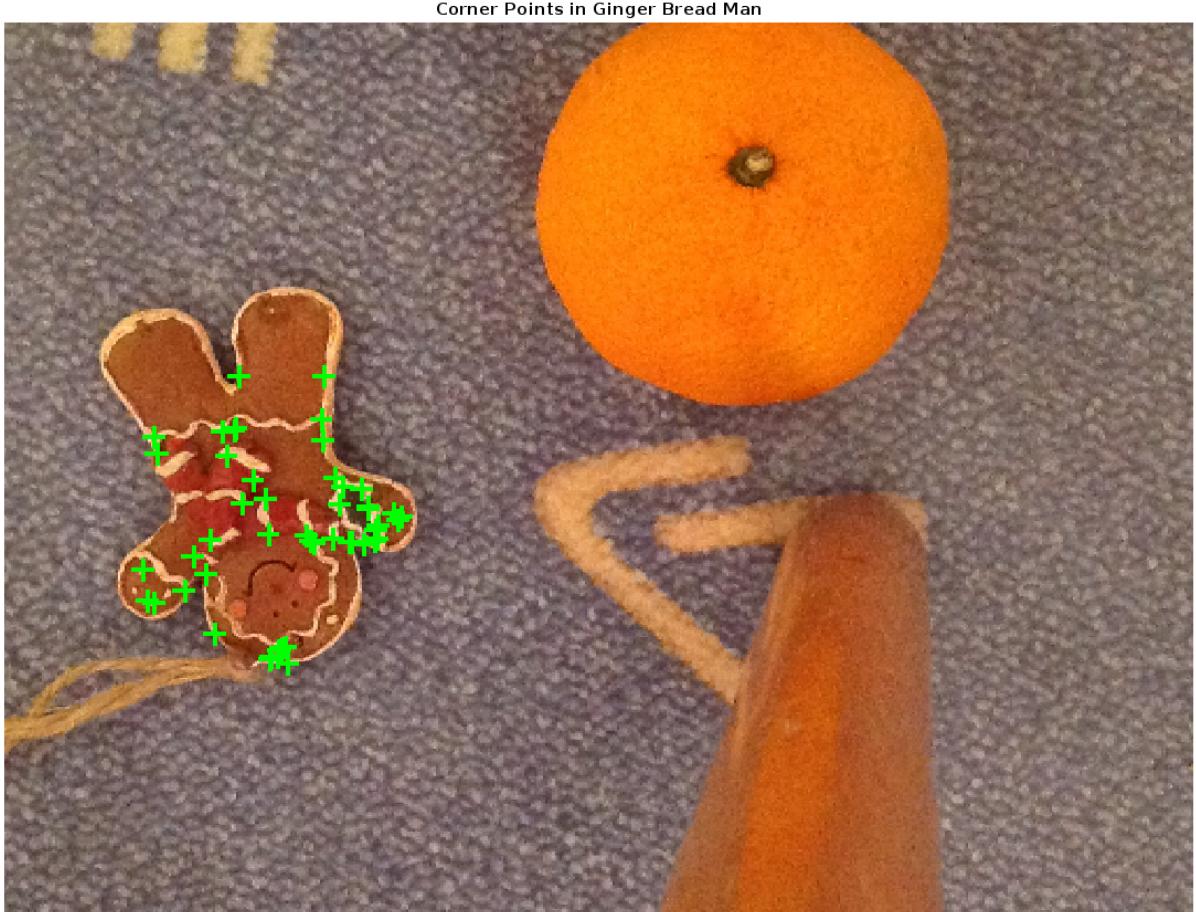


Figure 7: Corner points in the GingerBreadMan

We then proceed to save the optical flow vectors from the first image provided. This is then supplied to the second image for a better and updated determination of flow velocity vectors. As observed, the flow vectors centralise over the contours of the GingerBreadMan, indicating a successful flow estimation.



Figure 8: Flow vectors superimposed on the second image of GingerBreadMan

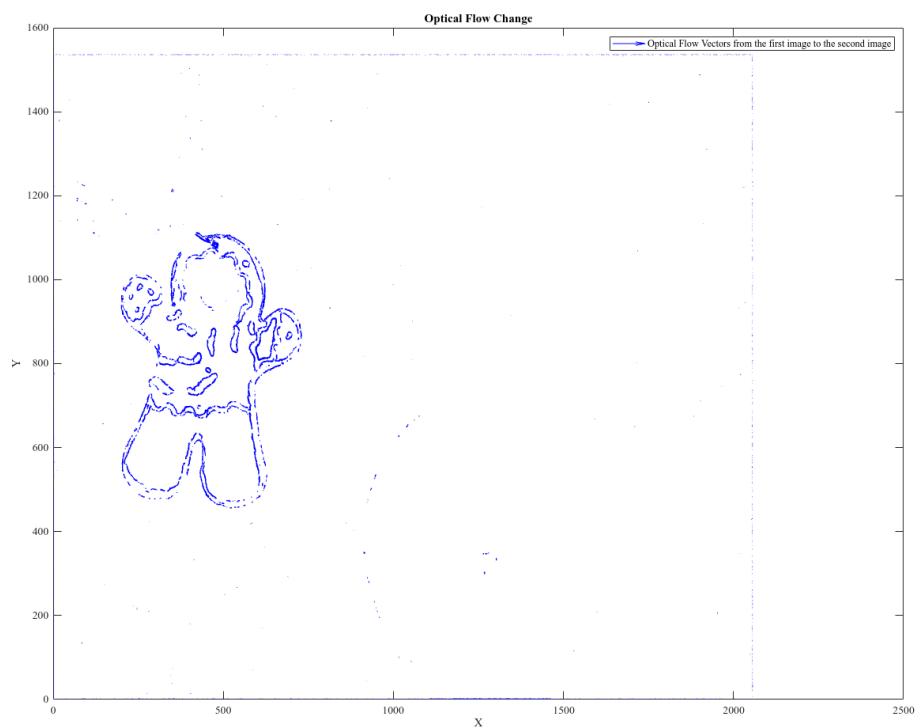


Figure 9: Plot of the flow vectors observed after Lucas-Kanade optical estimation

Conclusion

Here, we have tested the flow estimate for two videos, namely, the movement of a red square on a white background and a gingerbread cookie. Passing flow vectors from frame to frame has resulted in a better estimate of the next position of the object of interest, as seen in figure 6, as the initial estimate proved to contain high RMSE values at the start, but eventually converged.

The RMSE errors proved to be satisfactory in the case of the Red Square video and, indeed, were highly accurate as they matched the ground truth data very well. Additionally, the plot of the flow vectors from the gingerbreadman matched its contours, proving that the flow estimate is precise.

3 Task 3 - Automatic Detection of Moving Objects in a Sequence of Video Frames

3.1 Frame Differencing Approach

Background subtraction is an essential tool in Image Processing to determine the object of interest. A very naive approach is to use basic frame subtraction in order to extract the moving frames. This method is extremely computationally light and can produce reasonable results [4].

We proceed by considering the first frame of the video as the background and subtracting each subsequent frame from the former. This is done by binary thresholding the result with a customizable threshold to vary. Additionally, the frame of reference is updated for each step. The procedure can easily be cross-referenced by taking Task 3, part 1 of the Appendix, into account.

As observed, thresholding is very crucial to obtaining accurate results. Otherwise, insignificant frames get overlooked in frame difference due to increased frame sensitivity.



Figure 10: Frame Difference with a binary thresholding of 10 and 25, at frame 140

Analysis

The swaying of trees produces movable pixels and hence gets through the frame difference approach. However, the frame sensitivity can be reduced by varying the binary threshold. However, it may be difficult to capture subtle changes in the background while following this method. Moreover, there is no filtering step in frame differencing if noise is present in the video stream.

3.2 Gaussian Mixture Models

Gaussian Mixture Models are probabilistic models that represent small models that are probabilistic within an overall distribution. They break down complex multi-modal distributions into smaller Gaussian models in the aspects of mean and variance [5].

First, we supply an initial amount of frames to train the foreground detector. Hence, any additional training frames after a specific threshold do not affect the performance of the overall detector. A foreground mask is created after the model is trained using the Gaussian mixture models. This mask is then applied to each frame of the incoming video stream. Several tests can be made by varying the number of training frames and the number of Gaussian models by setting the initial parameters in T3 part2 of the Appendix.



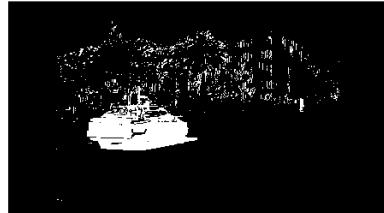
Foreground at Frame: 140 with 1 training frames & 2 Gaussians



Foreground at Frame: 140 with 1 training frames & 20 Gaussians



Foreground at Frame: 140 with 10 training frames & 2 Gaussians



Foreground at Frame: 140 with 10 training frames & 20 Gaussians

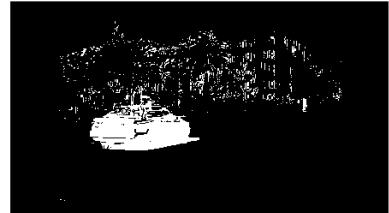


Figure 11: GMM with $n_frames = [1, 10]$, and $n_gaussians = [2, 20]$ at frame 140

Analysis

By analyzing the figure 11, we can observe that the most substantial parameter to vary is the number of training frames. As the training frames increase and saturate, the swaying of trees and the man walking in the background are observed as the filtered output. The Gaussian submodel split for 2 and 20 does not have a significant impact on the overall background subtraction.

4 Task 4 - Robot Treasure Hunting

4.1 Easy Case

We proceed with the underlying blueprint for all the cases as follows:

1. Converting the RGB image into a binary image
2. Analyzing the connected components in the image
3. Drawing and labelling the bounding boxes in the image

After getting the binary image, we label the connected clusters of pixels in an enumeration. Therefore, an index is assigned to every cluster of white pixels detected. A good way to visualize the detected clusters is by using the *label2rgb()* function. This is clearly seen in the figure 12, the neon pixel clusters are the labelled components [6]. All subsequent functions mentioned here can be referenced in Task 4 of the Appendix.

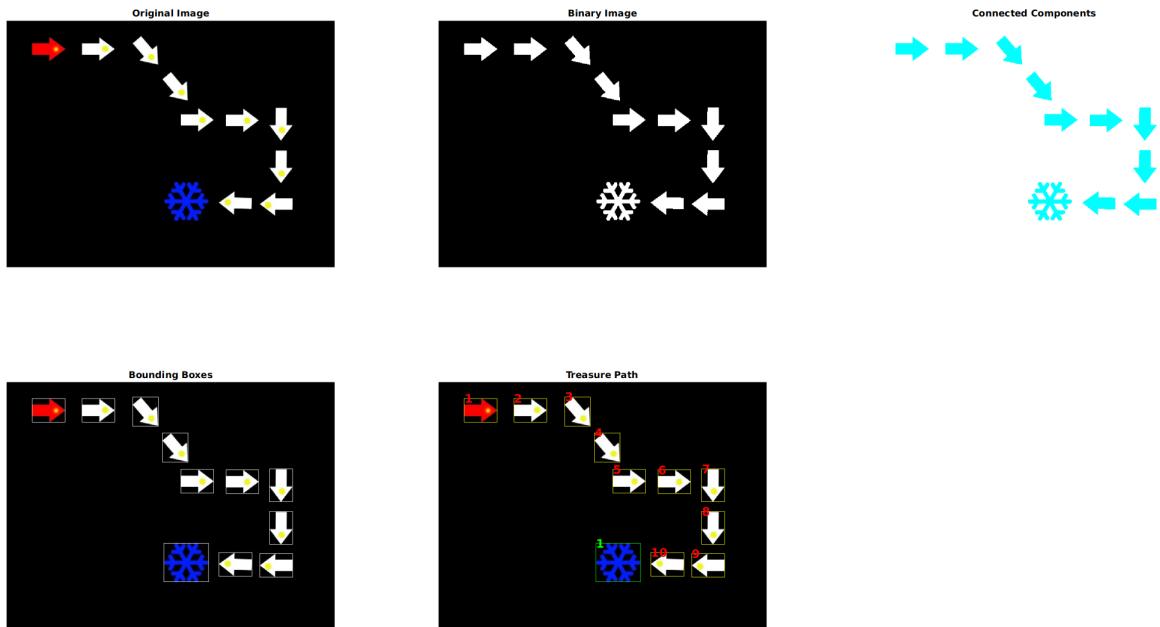


Figure 12: Intermediate steps in the Easy Case

This is done so we can get crisp and defined properties for the region properties. This helps us to get more accurate bounding boxes, centroids, and areas. After the bounding boxes are drawn and labelled, the main task is to label which boxes are arrows, treasures, or false detections. This is determined by the *arrow_finder()* function.

First, let us define the conditions for a bounding box to be an arrow:

1. The length and breadth of the bounding box need to be of a minimum length and breadth in order to be an effective detection.
2. The centroid of the bounding box has to be of white or yellow colour.

Then, let us define the suitable outlines for a bounding box to be a treasure:

1. The length and breadth of the bounding box need to be of a minimum length and breadth in order to be an effective detection.
2. The area of the bounding box is sufficiently larger than the rest of the boxes being detected.

Once we have extracted the indices of the treasure and arrows from the list of bounding boxes, the next task is to determine the starting point of the path to follow. This is easily determined as, in all of the three tasks, we only have a single red arrow as the starting point. Therefore, the starting index is determined by whether the centroid of the arrow bounding box is red.

Next, we need to determine the next arrow for traversal from the start. This is determined by the *next_object_finder()* function. This returns the next object in the path, whether it is an arrow or a treasure.

1. Effectively, the intuitive way to travel to the treasure is to follow from the tip of the current arrow to the base of the next arrow. This is done by calculating the distance from the current yellow centroid to the centroid of the next object.
2. The yellow centroid is computed in the *get_yellow_centroid()* function. This is done by cropping a region of interest and applying an HSV yellow mask to extract the bounding box of the yellow dot. The yellow centroid is then computed using region properties.
3. Next, we iterate through all the objects detected by *arrow_finder()* and omit objects already in the treasure path.
4. The centroid of the iterative object is computed, and the effective difference is monitored with the yellow centroid of the current arrow.

5. Once the index of the object is determined with the lowest effective distance, then the object's index is returned.
6. While iterating, the distance with the treasure is monitored, and a check for the area is done to prevent unwanted treasures.
7. After the index of the next object is determined, it gets appended in the path to follow. This is repeated until the object encounters a treasure.

We then proceed to draw yellow bounding boxes in the path if it is an arrow or green if it is a treasure. All the steps above can be inferred from the codebase in the appendix and the figure 12.

4.2 Medium Case

Here, the crux of this case is to determine which arrows are useful and which treasure is the destination. The earlier case for monitoring the distance between the current yellow centroid and the target object centroid still proves beneficial.

Additionally, proper thresholds for the treasure area need to be determined. This is easily found by monitoring the areas in the region properties of the bounding boxes detected. Hence, the path properly detects the fish instead of the iris flower.

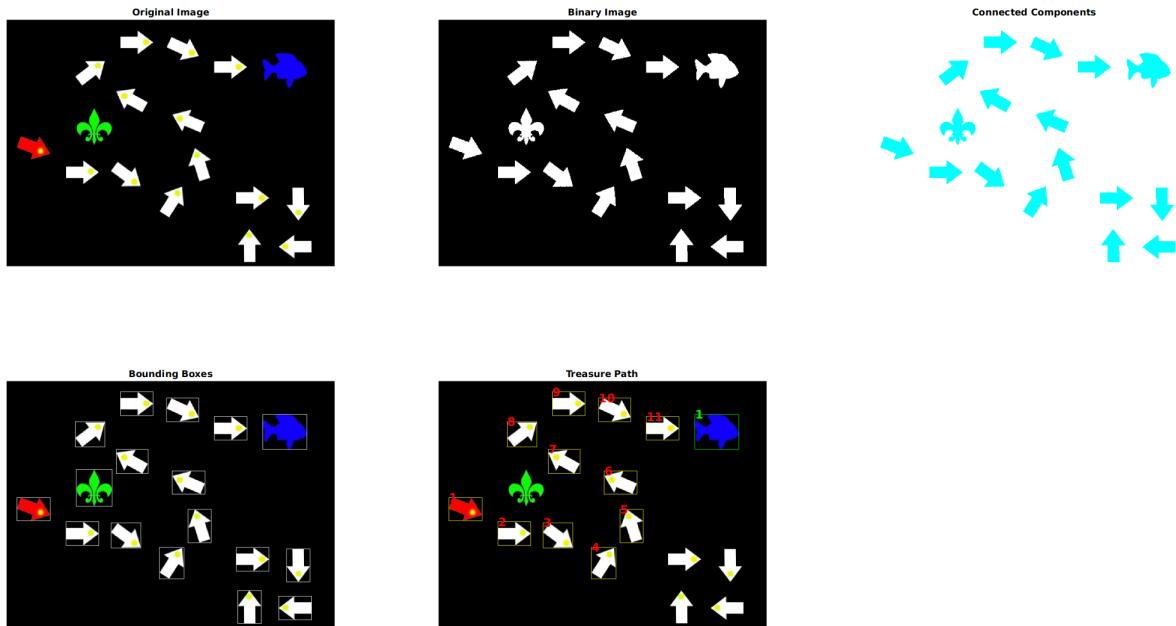


Figure 13: Intermediate steps in the Medium Case

The path does not choose the four arrows in the bottom right corner in the figure 13, as after the 5th arrow, the effective distance of the 6th labelled arrow is far smaller than the corresponding arrows.

$$effective_distance = \text{norm}(\text{current_yellow_centroid} - \text{target_centroid})$$

4.3 Hard Case

Following the approach in the medium case, it segregates the treasures of interest by simply taking the greatest pair of areas in the image. This works because the question mark box in figure 14 is smaller than the clover and the sun object.

However, since we have two treasures here, we need to consider the next point to follow to be from the intermediate treasure rather than start from an arrow. Here, the effective distance does not work directly, as there is no yellow centroid from which to calculate.

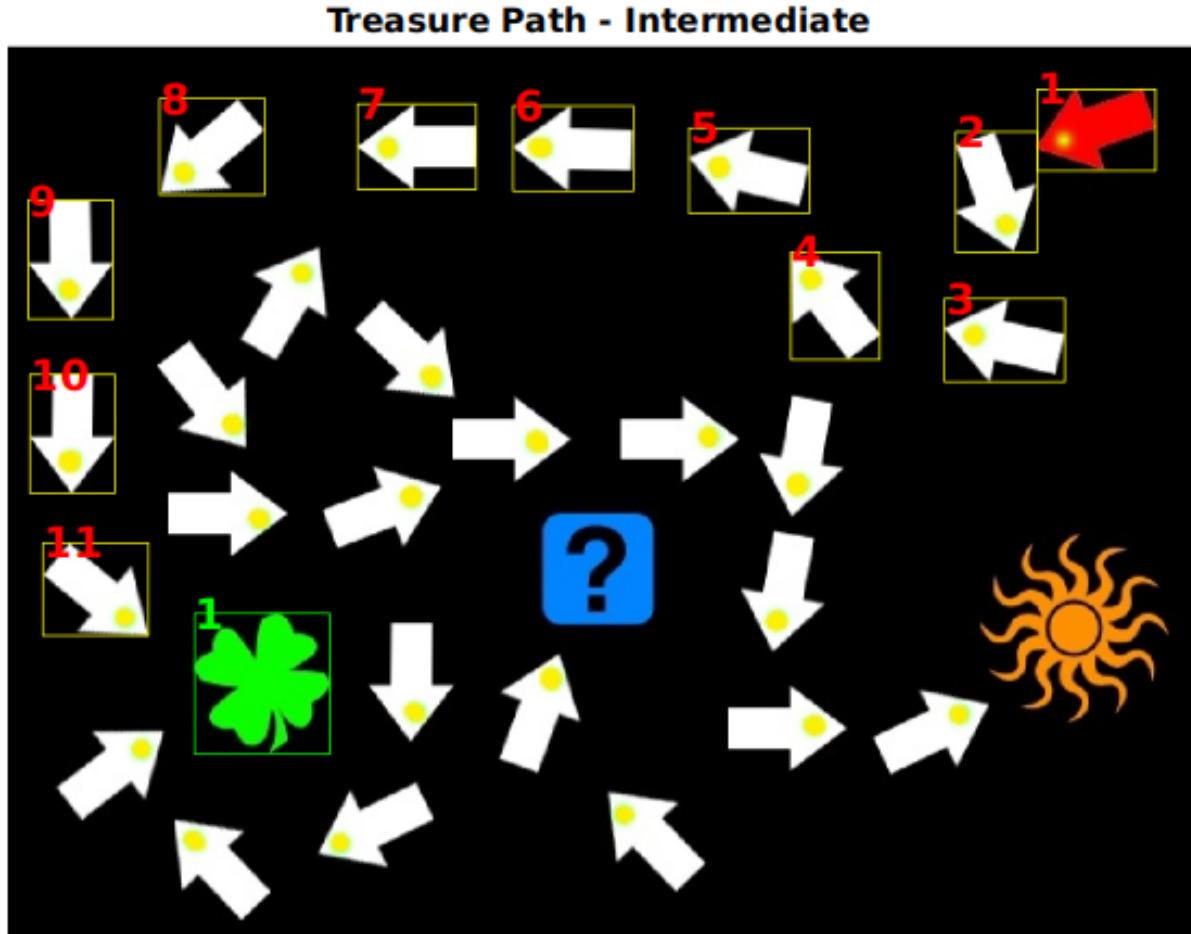


Figure 14: Intermediate treasure in Hard Case

Here, we skip all the next arrows lying in the width of the bounding box of the first clover treasure. This helps us to avoid "**jumps**" or "**diagonal movements**" while following the path to the second treasure. After following this simple trick, we can follow the effective distance defined in the easy and medium tasks, respectively.

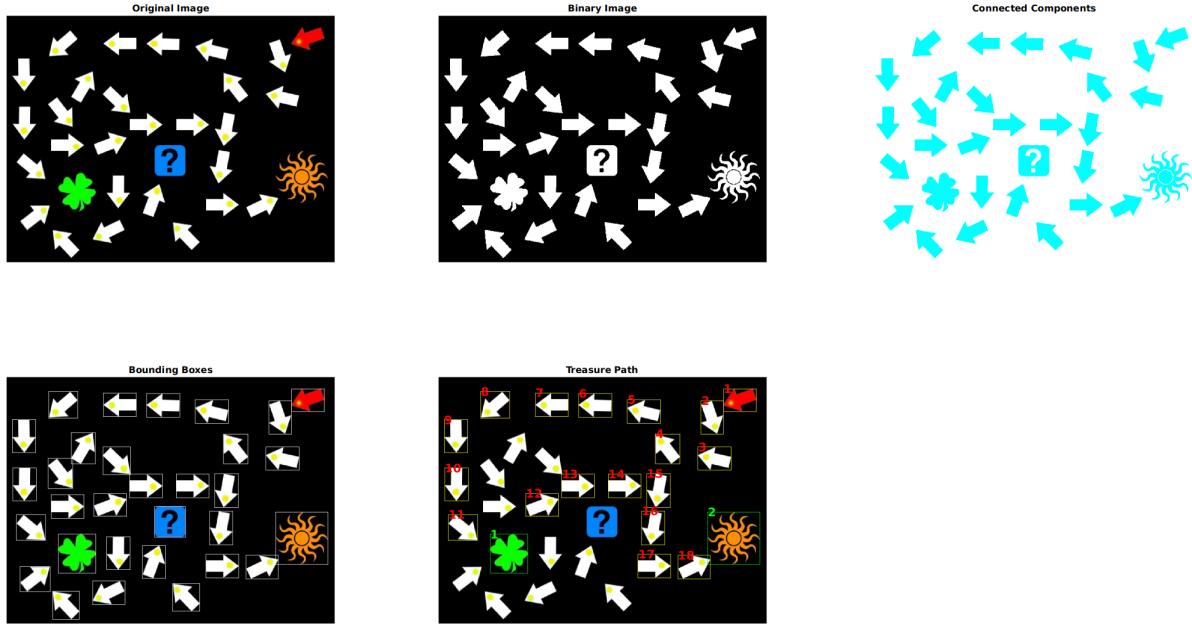


Figure 15: Intermediate steps in Hard Case

5 Task 5 - Convolutional Neural Networks for Image Classification

5.1 Accuracy of Default LeNet5

The default LeNet5 architecture is tested on the popular digits dataset, which consists of 10 classes. The layers and the training options have been pre-defined to match the architecture of Google's LeNet5 [7].

```
% Define the LeNet-5 architecture
layers = [
    imageInputLayer([image_size 1])

    convolution2dLayer(5,6)
    averagePooling2dLayer(2,'Stride',2)

    convolution2dLayer(5,16)
    averagePooling2dLayer(2,'Stride',2)

    fullyConnectedLayer(120)
    fullyConnectedLayer(84)

    fullyConnectedLayer(num_classes)
    softmaxLayer
    classificationLayer
];
```

As observed in the above code snippet, the architecture has two hidden layers and two fully connected layers. In addition, the training step uses a Stochastic Gradient with momentum as its optimizer. A custom change was introduced in the training options to use the GPU by default to speed up the model training. For more information on the GPU configuration, refer to Task 5 of the Appendix.

As seen from the figure 16, the accuracy of the model is about 63.26%.

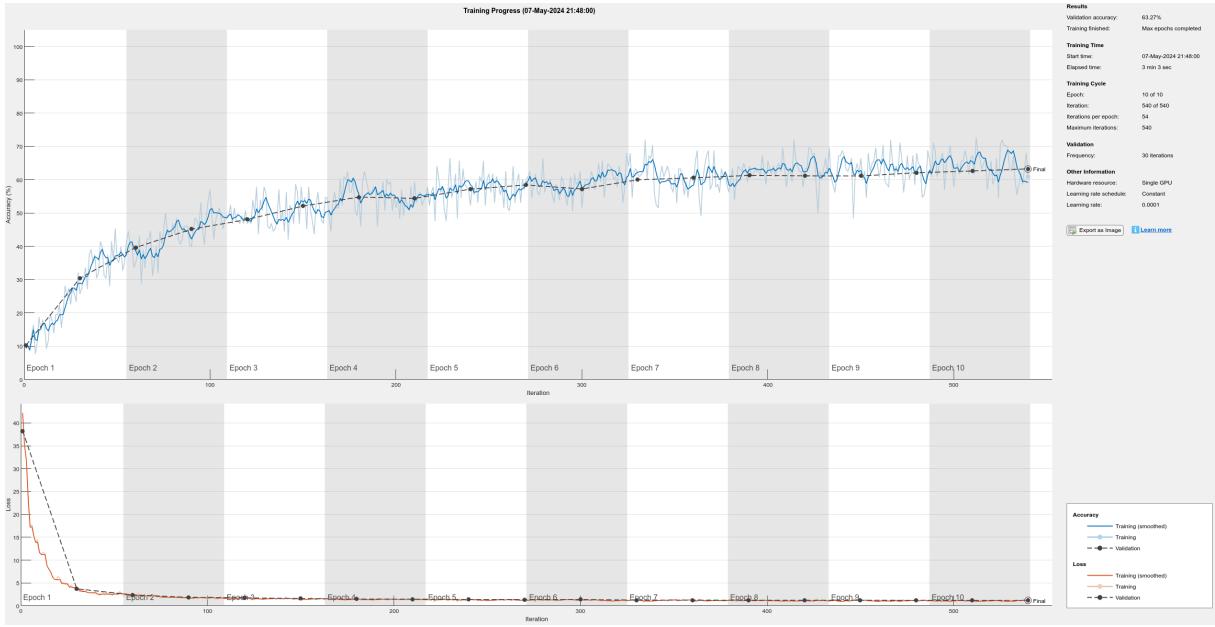


Figure 16: Training Progress on LeNet5

5.2 Precision, Recall and F1 Score of Default LeNet5 Model

In order to calculate the precision, recall and f1 score, we need to plot the confusion matrix from the training results.

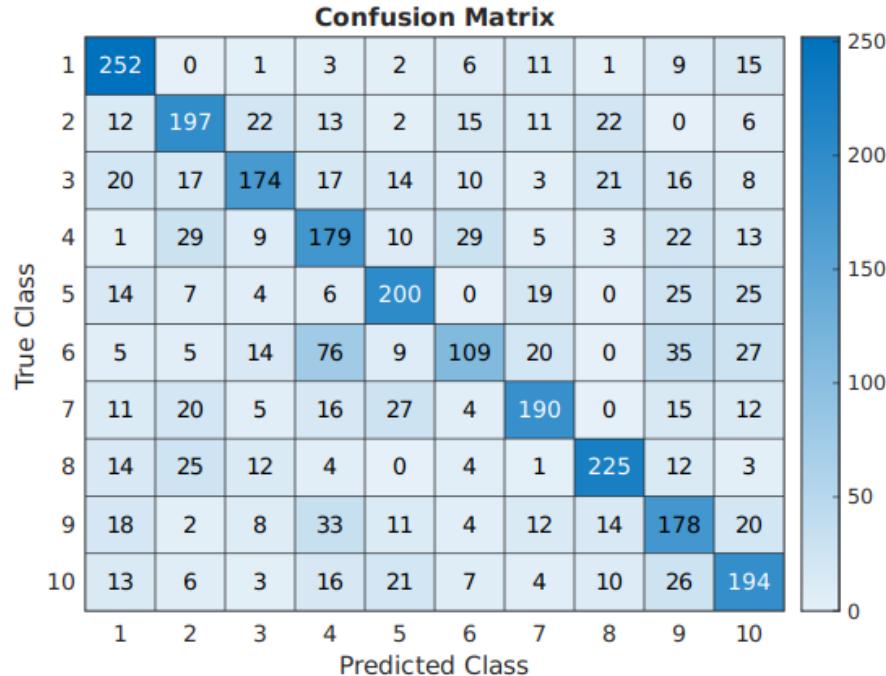


Figure 17: Confusion Matrix of the LeNet5 model

The following are the equations for calculating the precision, recall, and f1 score for each of the classes:

```
precision = diag(confusionMatrix) ./ sum(confusionMatrix, 2);  
recall = diag(confusionMatrix) ./ sum(confusionMatrix, 1)';  
f1Score = 2 * (precision .* recall) ./ (precision + recall);
```

Using the above equations, we get the following results:

$$precision = [0.84, 0.65, 0.58, 0.59, 0.66, 0.36, 0.63, 0.75, 0.59, 0.64]$$

$$recall = [0.7, 0.63, 0.69, 0.49, 0.67, 0.57, 0.68, 0.76, 0.52, 0.60]$$

$$f1_score = [0.76, 0.64, 0.63, 0.53, 0.67, 0.44, 0.65, 0.75, 0.55, 0.62]$$

$$overall_precision = 0.63$$

$$overall_recall = 0.63$$

$$overall_f1_score = 0.62$$

5.3 Improvement of the LeNet5 Model

The following methods have been employed in order to increase the overall accuracy of the system:

1. **Regularization methods:** L2 regularization has been tuned for maximum performance and applied to the fully connected layer. In addition, to prevent overfitting, a dropout method of a tuned probability measure is also added after the fully connected layers.
2. **Activation function:** Two ReLU layers have been applied to the two hidden layers of the custom architecture in order to provide an accurate activation variation to the output of each layer.
3. **Custom Architecture Modifications:** A couple of hidden layers have been applied prior to fully connected layers. In addition, a batch normalization layer has been applied to the output of every 2D convolution filter in order to offer a base unit output to the next layers.

4. Hyperparameter Tuning: The filter size and the number of filters for each hidden layer have been tested and modified for maximum performance. Moreover, the optimizer has been changed from *sgdm* to *adam*, because of its robustness.

The newly formed layers look as follows:

```
layers = [
    imageInputLayer([image_size 1])

    % Hidden Layer 1
    convolution2dLayer(filterSize, numFilters, 'Padding', 'same')
    batchNormalizationLayer
    reluLayer
    maxPooling2dLayer(2, 'Stride', 2)

    % Hidden Layer 2
    convolution2dLayer(filterSize, numFilters, 'Padding', 'same')
    batchNormalizationLayer
    reluLayer
    maxPooling2dLayer(2, 'Stride', 2)

    fullyConnectedLayer(fullyConnectedLayerSize, ...
        'WeightL2Factor', l2_reg, 'BiasL2Factor', l2_reg)
    dropoutLayer(0.2)

    fullyConnectedLayer(num_classes)
    softmaxLayer

    classificationLayer
];
```

Now, we proceed to train the custom model with the *adam* optimizer and analyse its output metrics as well.

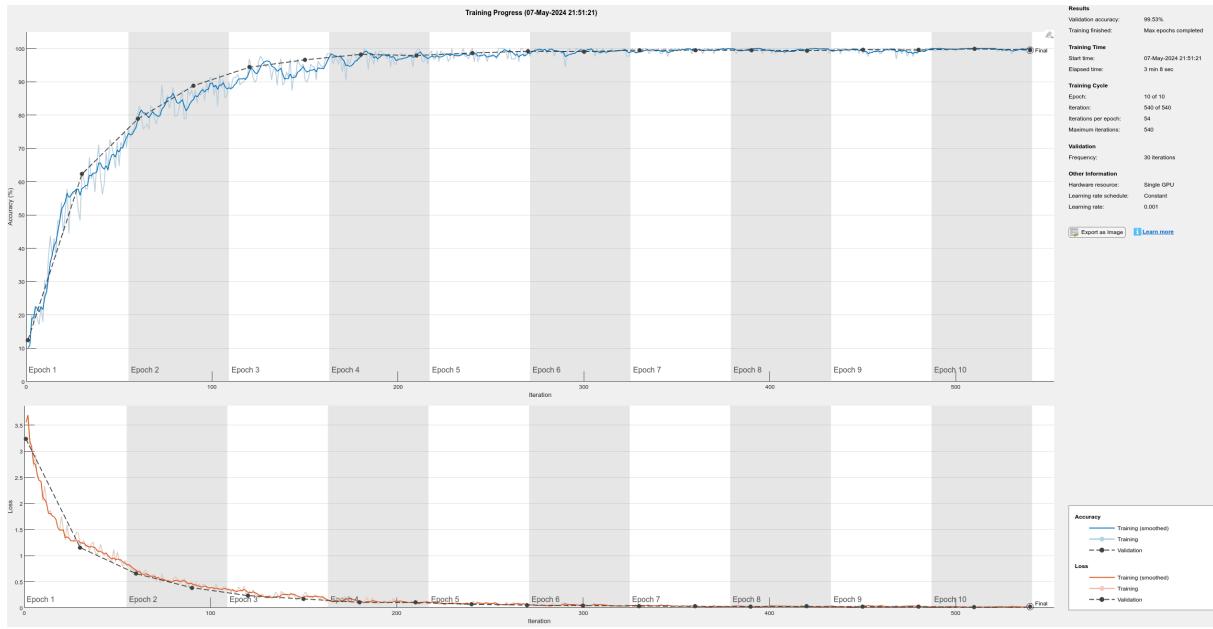


Figure 18: Training Progress on Custom model

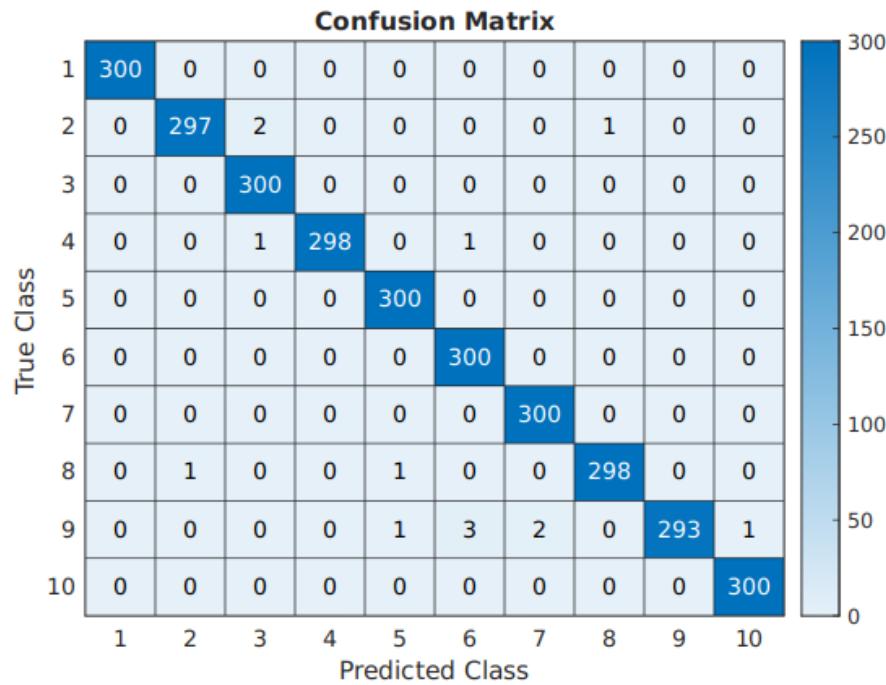


Figure 19: Confusion Matrix of the Custom model

The initial inference from both the training and loss and confusion plots is that the accuracy has bumped up to 99.5%, and all the diagonals of the confusion matrix have been populated.

$$precision = [1.00, 0.99, 1.00, 0.99, 1.00, 1.00, 1.00, 0.99, 0.97, 1.00]$$

$$recall = [1.0, 0.99, 0.99, 1.00, 0.99, 0.98, 0.99, 0.99, 1.00, 0.99]$$

$$f1_score = [1.00, 0.99, 0.99, 0.99, 0.99, 0.99, 0.99, 0.99, 0.98, 0.99]$$

$$overall_precision = 0.99$$

$$overall_recall = 0.99$$

$$overall_f1_score = 0.99$$

5.4 Ethics in Computer Vision

5.4.1 Positives of using Computer Vision

The following are some of the benefits of using any computer vision library in our day-to-day lives:

1. **Healthcare:** Early detection of cancers and tumours, especially through MRI imaging, can be diagnosed early if pre-trained custom models are used. This can benefit the facility and patients by lowering costs.
2. **Security:** Popular doorbell security systems such as Ring has already been deployed in many households to reduce intruder entry. Simple motion models and HaarCascade face detectors can be employed in CCTV systems to prevent theft and robbery.
3. **Automating Tasks:** Industry lines can be employed with custom models for picking and placing operations, labelling, identifying defects in manufacturing and so on. This can reduce mundane labour and can improve quality of life.

5.4.2 Challenges and Mitigation of using Computer Vision

The following are some of the downsides of using machine vision techniques in our way of life [8]:

1. **Scarcity of Labelled Data:** Certain use cases require a large number of labelled datasets, which can prove to be computationally heavy and expensive to manage.

Mitigation - Use of synthetic data generation methods such as changing the background, flipping and rotating the object of interest.

2. **Privacy:** Deepfakes have been recently introduced into social media platforms in great numbers due to AI image generation models such as Midjourney and DALLE. This can harm one's reputation and can prove to be hazardous to one's mental health.

Mitigation - The generation of these images must be limited to fictional people. Each image must be watermarked in advance, especially in the EXIF data, which is not apparently obvious.

3. **Greater number of dimensions in data:** It is hard to determine directly the leading behaviour in an outcome when presented with many features. This can lead to an unnecessary increase in training time and can prove to be cost-ineffective.

Mitigation - Dimensional reduction techniques such as Principal Component Analysis can be applied to find the most significant components of the output vector.

5.4.3 Equality and Diversity in Computer Vision

1. **Equality:** A popular issue in image processing is the facial bias in detection techniques. This can usually happen when using a biased dataset to train the model. This can result in the model detecting only some of the faces well while neglecting the accuracy of detecting the people outside the particular region.

2. **Diversity:** This arises from the equality measure itself. Using data from diverse regions can lead to a greater number of generated data that is presented equally without any bias.

6 Appendix

- Task-1, Part-1: Introduction to Image Processing

```
1 % Copyright (c) 2024 Leander Stephen D'Souza
2 % Program to plot rgb histogram of an image
3
4 % Call the function with an example images
5 plot_rgb_histogram(imread('duckMallardDrake.jpg'));
6 plot_rgb_histogram(imread('Dog.jpg'));
7
8
9 % Function to plot rgb histogram of an image
10 function plot_rgb_histogram(img)
11     % Get the size of the image
12     [~, ~, channels] = size(img);
13
14     % Check if the image is rgb
15     if channels ~= 3
16         disp('The image is not an rgb image');
17         return;
18     end
19
20     % Initialize the histogram
21     histogram = zeros(256, 3);
22
23     % Calculate the histogram using imhist
24     for k = 1:channels
25         histogram(:, k) = imhist(img(:, :, k));
26     end
27
28     % Display the image and histogram in a single figure
29     font_size = 12;
30     figure;
31     subplot(1, 2, 1);
32     imshow(img);
33     title('Original Image');
34
35     subplot(1, 2, 2);
36     bar(histogram(:, 1), 'r');
37     hold on;
38     bar(histogram(:, 2), 'g');
39     bar(histogram(:, 3), 'b');
40     title('RGB Histogram','FontSize',font_size);
41     xlabel('Intensity','FontSize',font_size);
42     ylabel('Frequency','FontSize',font_size);
43     legend('Red', 'Green', 'Blue', 'Location', 'best');
44 end
```

- Task-1, Part-2: Edge Detection and Segmentation of Static Objects

```

1 % Copyright (c) 2024 Leander Stephen D'Souza
2
3 % Program to test Edge Detection and Segmentation of Static Objects
4
5
6 og_image = imread('duckMallardDrake.jpg');
7
8 % Edge Detection
9 edge_detection_analysis(og_image);
10
11
12 function edge_detection_analysis(img)
13     % 1. Apply Sobel Edge Detection
14     sobel_edge_image = edge(rgb2gray(img), 'sobel', 0.1);
15
16     % 2. Apply Canny Edge Detection
17     canny_edge_image = edge(rgb2gray(img), 'canny', 0.15);
18
19     % 3. Apply Prewitt Edge Detection
20     prewitt_edge_image = edge(rgb2gray(img), 'prewitt', 0.05);
21
22     % Plot all the images with appropriate titles
23     figure;
24     subplot(2, 2, 1);
25     imshow(img);
26     title('Original Image');
27
28     subplot(2, 2, 2);
29     imshow(sobel_edge_image);
30     title('Sobel Edge Detection');
31
32     subplot(2, 2, 3);
33     imshow(canny_edge_image);
34     title('Canny Edge Detection');
35
36     subplot(2, 2, 4);
37     imshow(prewitt_edge_image);
38     title('Prewitt Edge Detection');
39 end

```

- Task-2: Object Motion Detection and Tracking

```

1 % Copyright (c) 2024 Leander Stephen D'Souza
2
3 % Program to analyze Optical Flow in a Video Sequence
4 clc; clear; close all;
5
6 % Task 1: Find Corner Points
7 % find_corner_points(imread('red_square_static.jpg'), 'Red Square');
8 find_corner_points(imread('GingerBreadMan_first.jpg'), 'Ginger Bread Man');
9
10 % % Task 2: Optical Flow
11 first_image = imread('GingerBreadMan_first.jpg');
12 second_image = imread('GingerBreadMan_second.jpg');
13 estimate_optical_flow(first_image, second_image);
14
15 % % Task 3: Optical Flow in Video Sequence
16 video = VideoReader('red_square_video.mp4');
17 ground_truth = load('new_red_square_gt.mat');
18 % optical_flow_video(video, ground_truth);
19
20
21 % Function to analyze the optical flow in a video sequence
22 function optical_flow_video(video, ground_truth)
23     % Read the video frames
24     video_frames = read(video);
25     num_frames = video.NumFrames;
26
27     % Create an optical flow object
28     opticalFlow = opticalFlowLK('NoiseThreshold', 0.009);
29
30     % Load the actual corner trajectory
31     actual_trajectory = cell2mat(struct2cell(ground_truth));
32
33     % Save all the estimated trajectory of the corner points
34     predicted_trajectory = zeros(num_frames - 1, 2);
35
36     % Initialize the corner points
37     first_frame = rgb2gray(video_frames(:, :, :, 1));
38     corners = corner(first_frame, 'Harris');
39     corner_x = corners(1, 1);
40     corner_y = corners(1, 2);
41
42     % Initialize flow
43     flow = estimateFlow(opticalFlow, first_frame);
44
45     % Initialize corner trajectory
46     vx = flow.Vx(round(corner_y), round(corner_x));
47     vy = flow.Vy(round(corner_y), round(corner_x));
48

```

```

49 % Compute the new position by adding the velocity vector to the current
50 % position
51 x_new = corner_x + vx;
52 y_new = corner_y + vy;
53 predicted_trajectory(1, :) = [x_new, y_new];
54
55 % Iterate through the video frames
56 for i = 2:num_frames
57     % Read the current frame
58     current_frame = rgb2gray(video_frames(:,:,:,:i));
59
60     % Detect corners in the current frame
61     corners = corner(current_frame, 'Harris');
62
63     % Find the nearest corner to the previous corner
64     min_dist = 1000000;
65     for j = 1:size(corners, 1)
66         dist = sqrt((corners(j, 1) - x_new)^2 + (corners(j, 2) - y_new)^2);
67         if dist < min_dist
68             min_dist = dist;
69             nearest_corner = corners(j, :);
70         end
71     end
72
73     % Update the corner position
74     corner_x = nearest_corner(1);
75     corner_y = nearest_corner(2);
76
77     % Estimate the optical flow for the current frame
78     flow = estimateFlow(opticalFlow, current_frame);
79
80     % Compute the new position by adding the velocity vector to the current
81     % position
82     vx = flow.Vx(round(corner_y), round(corner_x));
83     vy = flow.Vy(round(corner_y), round(corner_x));
84
85     % Update the track with the new position
86     predicted_trajectory(i, :) = [corner_x + vx, corner_y + vy];
87 end
88
89 % Plot the estimated corner trajectory and the actual corner trajectory
90 last_frame = video_frames(:,:,:,:num_frames);
91
92 imshow(last_frame);
93 hold on
94 plot(predicted_trajectory(:, 1), predicted_trajectory(:, 2), 'r');
95 plot(actual_trajectory(:, 1), actual_trajectory(:, 2), 'b');
96 title('Estimated vs Actual Corner Trajectory');
97 xlabel('X');

```

```

96     ylabel('Y');
97     legend('Estimated Trajectory', 'Actual Trajectory');
98     hold off
99
100    % Get RMSEx, RMSEy, and RMSE
101    for i = 1:num_frames-1
102        RMSE_x(i) = sqrt((actual_trajectory(i+1, 1) - predicted_trajectory(i, 1)) ^ 2);
103        RMSE_y(i) = sqrt((actual_trajectory(i+1, 2) - predicted_trajectory(i, 2)) ^ 2);
104        RMSE(i) = sqrt((actual_trajectory(i+1, 1) - predicted_trajectory(i, 1)) ^ 2 ...
105                      + ... ...
106                      (actual_trajectory(i+1, 2) - predicted_trajectory(i, 2)) ^ 2);
107    end
108
109    % Plot the RMSE values
110    figure;
111    plot(2:num_frames, RMSE_x, 'b');
112    hold on;
113    plot(2:num_frames, RMSE_y, 'r');
114    plot(2:num_frames, RMSE, 'g');
115    title('RMSE Values');
116    xlabel('Frame Number');
117    ylabel('RMSE');
118    legend('RMSE_x', 'RMSE_y', 'RMSE');
119    hold off;
120
121    % Print the RMSE values
122    fprintf('RMSE_x: %f\n', mean(RMSE_x));
123    fprintf('RMSE_y: %f\n', mean(RMSE_y));
124    fprintf('RMSE: %f\n', mean(RMSE));
125
126
127    % Function to estimate the optical flow between two images
128    function estimate_optical_flow(first_image, second_image)
129        % Convert the images to grayscale
130        first_image_gray = rgb2gray(first_image);
131        second_image_gray = rgb2gray(second_image);
132
133        % Find the optical flow of the pixels between the two images
134        opticalFlow = opticalFlowLK('NoiseThreshold', 0.01);
135
136        % The object stores the vectors from the first image to the second image
137        estimateFlow(opticalFlow, first_image_gray);
138        flow = estimateFlow(opticalFlow, second_image_gray);
139
140        % Plot the optical flow on the second image

```

```

141 figure;
142 imshow(second_image);
143 hold on;
144 plot(flow, 'DecimationFactor', [5 5], 'ScaleFactor', 15);
145 title('Optical Flow between the two images');
146 hold off;
147
148 % Plot the optical flow vectors
149 figure;
150 plot(flow);
151 xlabel('X');
152 ylabel('Y');
153 legend('Optical Flow Vectors from the first image to the second image');
154 title('Optical Flow Change');
155 end
156
157
158 % Function to find the corner points in an image
159 function find_corner_points(img, title_name)
160     % Convert the image to grayscale
161     img_gray = rgb2gray(img);
162
163     % Find the corner points in the image
164     corners = detectHarrisFeatures(img_gray);
165
166     % Save the strongest corner points
167     strongest_corners = corners.selectStrongest(50);
168
169     % Plot the image with the corner points
170     figure;
171     imshow(img);
172     hold on;
173     plot(strongest_corners.Location(:, 1), strongest_corners.Location(:, 2), 'g+', ...
174          'MarkerSize', 15, 'LineWidth', 3);
175     title("Corner Points in " + title_name);
176     hold off;
177 end

```

- Task-3, Part-1: Frame Differencing Approach

```

1 % Copyright (c) 2024 Leander Stephen D'Souza
2
3 % Program to analyze background subtraction using frame difference
4
5 close all; clear; clc; % clear the workspace
6
7 % read the video
8 source = VideoReader('car-tracking.mp4');
9
10 % set the threshold
11 thresh_arr = [10, 25];
12 n_cols = 2;
13 n_rows = 1 + length(thresh_arr) / n_cols;
14
15 figure
16 set(gcf, 'WindowState', 'maximized');
17
18 for i = 1:length(thresh_arr)
19     % call the function to subtract the background
20     source.CurrentTime = 0; % reset the video to the beginning
21     frame_difference(source, thresh_arr(i), n_rows, n_cols, i); % call the function
22         to subtract the background
23 end
24
25 % function to subtract the background
26 function frame_difference(source, thresh, n_rows, n_cols, i)
27     % read the first frame of the video as a background model
28     bg = readFrame(source);
29     bg_bw = rgb2gray(bg); % convert background to greyscale
30
31     frame_counter = 0; % initialize the frame counter
32     frame_break = 140;
33
34     % ----- process frames -----
35     % loop all the frames
36     while hasFrame(source)
37         frame_counter = frame_counter + 1; % increment the frame counter
38
39         if frame_counter > frame_break
40             break;
41         end
42
43         fr = readFrame(source); % read in frame
44         fr_bw = rgb2gray(fr); % convert frame to grayscale
45         fr_diff = abs(double(fr_bw) - double(bg_bw)); % cast operands as double to
46             avoid negative overflow

```

```

47      % if fr_diff > thresh pixel in foreground
48      fg = uint8(zeros(size(bg_bw)));
49      fg(fr_diff > thresh) = 255;
50
51      % update the background model
52      bg_bw = fr_bw;
53
54      % visualise the results and label the frame number with titles
55      subplot(n_rows, n_cols, 1), imshow(fr), title('Original Frame')
56      subplot(n_rows, n_cols, i + 2), imshow(fg), title('Foreground Pixels at
Frame: ' + string(frame_counter) + ' with threshold: ' + string(thresh))
57      drawnow;
58
59 end

```

- Task-3, Part-2: Gaussian Mixture Models

```

1 % Copyright (c) 2024 Leander Stephen D'Souza
2
3 % Program to analyze background subtraction using Gaussian mixture models
4
5 close all; clear; clc; % clear the workspace
6
7 % read the video
8 source = VideoReader('car-tracking.mp4');
9
10 % variable parameters
11 n_frames = [1, 10];
12 n_gaussians = [2, 20];
13
14 n_rows = length(n_frames);
15 n_cols = length(n_gaussians);
16
17 figure
18 set(gcf,'WindowState','maximized');
19
20 for i = 1:n_rows
21     for j = 1:n_cols
22         % call the function to subtract the background
23         source.CurrentTime = 0; % reset the video to the beginning
24         gaussian_mixture_models(source, n_frames(i), n_gaussians(j), n_rows + 1,
25         n_cols, i, j); % call the function to subtract the background
26     end
27 end
28
29 % Function to implement Gaussian Mixture Models
30 function gaussian_mixture_models(source, n_frames, n_gaussians, n_rows, n_cols, i,
31 j)
32 frame_counter = 0; % initialize the frame counter
33 frame_break = 140;
34
35 detector = vision.ForegroundDetector('NumTrainingFrames', n_frames,
36 'NumGaussians', n_gaussians);
37 % make new figure window
38
39 % ----- process frames -----
40 % loop all the frames
41 while hasFrame(source)
42     frame_counter = frame_counter + 1; % increment the frame counter
43
44     if frame_counter > frame_break
45         break;
46     end
47
48
49
50
51
52
53
54
55
56
57
58
59
60
61
62
63
64
65
66
67
68
69
70
71
72
73
74
75
76
77
78
79
80
81
82
83
84
85
86
87
88
89
90
91
92
93
94
95
96
97
98
99
100
101
102
103
104
105
106
107
108
109
110
111
112
113
114
115
116
117
118
119
120
121
122
123
124
125
126
127
128
129
130
131
132
133
134
135
136
137
138
139
140
141
142
143
144
145

```

```

46     fr = readFrame(source);      % read in frame
47
48     fgMask = step(detector, fr);    % compute the foreground mask by Gaussian
        mixture models
49
50     % create frame with foreground detection
51     fg = uint8(zeros(size(fr, 1), size(fr, 2)));
52     fg(fgMask) = 255;
53
54     % visualise the results after labelling the frame number with titles
55     subplot(n_rows, n_cols, 1), imshow(fr), title('Original Frame')
56     subplot(n_rows, n_cols, i * n_cols + j), imshow(fg), title(['Foreground at
        Frame: ' + string(frame_counter) + ' with ' + string(n_frames) + ' training
        frames & ' + string(n_gaussians) + ' Gaussians'])
57
58     drawnow
59
60 end

```

- Task-4: Robot Treasure Hunting

```

1 % Copyright (c) 2024 Leander Stephen D'Souza
2
3 % Program to implement Treasure Hunting using Arrow Detection
4 clear; clc; close all;
5
6 %% Treasure Hunting
7 treasure_hunter(imread('Treasure_easy.jpg'));
8 treasure_hunter(imread('Treasure_medium.jpg'));
9 treasure_hunter(imread('Treasure_hard.jpg'));
10
11
12 % Function to implement Treasure Hunting
13 function treasure_hunter(img)
14     figure;
15
16     % Show image
17     subplot(2, 3, 1);
18     imshow(img);
19     title('Original Image');
20
21     % Binarisation
22     bin_img = imbinarize(rgb2gray(img), 0.1);
23     subplot(2, 3, 2);
24     imshow(bin_img);
25     title('Binary Image');
26
27
28     % Extracting connected components
29     con_com = logical(bwlabel(bin_img));
30     subplot(2, 3, 3);
31     imshow(label2rgb(con_com));
32     title('Connected Components');
33
34     % Computing objects properties
35     props = regionprops(con_com);
36
37     % Drawing bounding boxes
38     n_objects = numel(props);
39     subplot(2, 3, 4);
40     imshow(img);
41     hold on;
42     for object_id = 1 : n_objects
43         rectangle('Position', props(object_id).BoundingBox, 'EdgeColor', 'w');
44     end
45     hold off;
46     title('Bounding Boxes');
47
48

```

```

49 % Finding arrows
50 [arrow_list, treasure_list] = arrow_finder(props, img);
51 all_boxes_list = [arrow_list, treasure_list];
52
53 % Finding red arrow (start point)
54 n_arrows = numel(arrow_list);
55 start_arrow_id = 0;
56 % check each arrow until find the red one
57 for arrow_num = 1 : n_arrows
58     object_id = arrow_list(arrow_num);      % determine the arrow id
59
60     % extract colour of the centroid point of the current arrow
61     centroid_colour = img(round(props(object_id).Centroid(2)), round(props(
62     object_id).Centroid(1)), :);
63
64     red = centroid_colour(:, :, 1);
65     green = centroid_colour(:, :, 2);
66     blue = centroid_colour(:, :, 3);
67
68     if red > 240 && green < 10 && blue < 10
69     % the centroid point is red, memorise its id and break the loop
70         start_arrow_id = object_id;
71         break;
72     end
73 end
74
75 % Hunting
76 cur_object = start_arrow_id; % start from the red arrow
77 path = cur_object; % memorise the path
78
79 for treasure_idx = 1 : numel(treasure_list)
80     % while the current object is an arrow, continue to search
81     while ismember(cur_object, arrow_list)
82         % find the next closest arrow
83         prev_object = cur_object;
84         cur_object = next_object_finder(cur_object, props, all_boxes_list, img,
85         path);
86         path(end + 1) = cur_object;
87     end
88     % Found treasure
89     cur_object = prev_object;
90     break;
91 end
92
93 % helper variable for drawing the treasure
94 treasure_counter = 0;
95 treasure_found = false;
96 treasure_offset = 0;

```

```

96 % visualisation of the path
97 subplot(2, 3, 5);
98 imshow(img);
99 hold on;
100
101 for path_element = 1 : numel(path)
102     object_id = path(path_element); % determine the object id
103
104     % reset for treasure
105     if ismember(object_id, treasure_list)
106         treasure_counter = treasure_counter + 1;
107         treasure_found = true;
108         rectangle('Position', props(object_id).BoundingBox, 'EdgeColor', 'g');
109         str = num2str(treasure_counter);
110         text(props(object_id).BoundingBox(1), props(object_id).BoundingBox(2),
111             str, 'Color', 'g', 'FontWeight', 'bold', 'FontSize', 14);
112         continue;
113     end
114
115     if treasure_found
116         treasure_found = false;
117         treasure_offset = 1;
118         continue;
119     end
120
121     rectangle('Position', props(object_id).BoundingBox, 'EdgeColor', 'y');
122     str = num2str(path_element - treasure_counter - treasure_offset);
123     text(props(object_id).BoundingBox(1), props(object_id).BoundingBox(2), str,
124         'Color', 'r', 'FontWeight', 'bold', 'FontSize', 14);
125     end
126
127 hold off;
128 title('Treasure Path - Intermediate');
129
130 % Function to find the arrow objects
131 function [arrow_list, treasure_list] = arrow_finder(props, img)
132     n_objects = numel(props);
133     arrow_list = zeros(1, n_objects);
134     treasure_list = zeros(1, n_objects);
135     min_length = 5;
136     treasure_min_area = 3014;
137     arrow_count = 0;
138     treasure_count = 0;
139
140     % iterate over all objects
141     for object_id = 1 : n_objects
142         % extract the bounding box of the current object
143         bbox = props(object_id).BoundingBox;

```

```

143
144     % extract the length and breadth of the bounding box
145     length = bbox(3);
146     breadth = bbox(4);
147
148     % check if the object is of sufficient size
149     if breadth > min_length && length > min_length
150         % extract the centroid of the current object, if the colour of the
151         % centroid is white, then it is an arrow
152         centroid_colour = img(round(props(object_id).Centroid(2)), round(props(
153             object_id).Centroid(1)), :);
154
155         red = centroid_colour(:, :, 1);
156         green = centroid_colour(:, :, 2);
157         blue = centroid_colour(:, :, 3);
158
159         % Use the fact that the arrow is red or white
160         if (red > 240 && green > 240 && blue > 240) || (red > 240 && green < 10
161             && blue < 10)
162             arrow_count = arrow_count + 1;
163             arrow_list(arrow_count) = object_id;
164         elseif props(object_id).Area > treasure_min_area
165             treasure_count = treasure_count + 1;
166             treasure_list(treasure_count) = object_id;
167         end
168     end
169
170 end
171
172 % Function to find the next object in the path
173 function [cur_object] = next_object_finder(cur_object, props, all_boxes_list, img,
174     path)
175     % get the bounding box of the current object
176     bbox_cur_object = props(cur_object).BoundingBox;
177
178     % calculate the current yellow centroid
179     cur_yellow_centroid = get_yellow_centroid(img, bbox_cur_object);
180
181     % find the closest object to the current object
182     min_dist = inf;
183     critical_dist = 100;
184     treasure_area = 3014;
185     cur_object = 0;
186
187     for object_idx = 1 : numel(all_boxes_list)
188         object = all_boxes_list(object_idx);

```

```

188 % check if the object is not in the path
189 if ~ismember(object, path)
190     % get the centroid of the object
191     object_centroid = props(object).Centroid;
192
193     % get the distance between the current object and the object
194     dist = norm(cur_yellow_centroid - object_centroid);
195
196     % get area
197     if props(object).Area > treasure_area
198         % monitor intermediate treasure distance
199         if dist < critical_dist
200             critical_dist = dist;
201             cur_object = object;
202         end
203     end
204
205     % check if the object is the closest object
206     if dist < min_dist
207         min_dist = dist;
208         cur_object = object;
209     end
210 end
211
212 end
213
214 % Function to extract yellow centroid from a bounding box
215 function yellow_centroid = get_yellow_centroid(img, bbox)
216     % extract the bounding box of the current object
217     bbox = round(bbox);
218     x = bbox(1);
219     y = bbox(2);
220     width = bbox(3);
221     height = bbox(4);
222     visualize = false;
223
224     % extract the region of interest
225     roi = img(y : y + height, x : x + width, :);
226
227     % convert the region of interest to the HSV colour space
228     hsv_roi = rgb2hsv(roi);
229
230     % extract the hue, saturation and value channels
231     hue = hsv_roi(:, :, 1);
232     saturation = hsv_roi(:, :, 2);
233     value = hsv_roi(:, :, 3);
234
235     % threshold the hue channel to extract the yellow pixels
236     yellow_mask = hue > 0.1 & hue < 0.2;

```

```

237
238 % threshold the saturation channel to extract the yellow pixels
239 yellow_mask = yellow_mask & saturation > 0.5;
240
241 % threshold the value channel to extract the yellow pixels
242 yellow_mask = yellow_mask & value > 0.5;
243
244 % check if the yellow mask is empty
245 if ~any(yellow_mask(:))
246     yellow_centroid = [0, 0];
247 else
248     % find the centroid of the yellow pixels
249     yellow_centroid = regionprops(yellow_mask, 'Centroid');
250     yellow_centroid = yellow_centroid.Centroid;
251     yellow_centroid = [yellow_centroid(1) + x, yellow_centroid(2) + y];
252
253 % visualisation
254 if visualize
255     imshow(img);
256     hold on;
257     rectangle('Position', bbox, 'EdgeColor', 'r');
258     plot(yellow_centroid(1), yellow_centroid(2), 'ro');
259     hold off;
260     pause(0.5);
261 end
262 end
263 end

```

- Task-5: Convolutional Neural Networks for Image Classification

```

1 % Copyright (c) 2024 Leander Stephen D'Souza
2
3 % Program to implement Google LeNet CNN for Image Classification for the Digit
4 % Dataset
5
6 % define constants
7 image_size = [32, 32];
8 num_classes = 10;
9
10 % Load the digit dataset
11 digitDatasetPath = fullfile(toolboxdir('nnet'), 'nndemos', ...
12     'nndatasets', 'DigitDataset');
13 %
14 imds = imageDatastore(digitDatasetPath, ...
15     'IncludeSubfolders', true, 'LabelSource', 'foldernames');
16
17 imds.ReadFcn = @(loc)imresize(imread(loc), image_size);
18
19 % Split the data into training and validation datasets
20 [imdsTrain, imdsValidation] = splitEachLabel(imds, 0.7, 'randomized');
21
22 % Define the LeNet-5 architecture
23 [lenet_layers, lenet_options] = lenet_architecture(image_size, num_classes,
24     imdsValidation);
25
26 % Define custom architecture
27 [custom_layers, custom_options] = custom_architecture(image_size, num_classes,
28     imdsValidation);
29
30 for i = 1:2
31     if i == 1
32         layers = lenet_layers;
33         options = lenet_options;
34     else
35         layers = custom_layers;
36         options = custom_options;
37     end
38
39 % Train the network
40 net = trainNetwork(imdsTrain, layers, options);
41
42 % Classify validation images and compute accuracy
43 YPred = classify(net, imdsValidation);
44 YValidation = imdsValidation.Labels;
45 accuracy = sum(YPred == YValidation) / numel(YValidation);
46 fprintf('Accuracy of the network on the validation images: %f\n', accuracy);

```

```

46
47 % Calculate the confusion matrix
48 confusionMatrix = confusionmat(YValidation, YPred);
49
50 % Display the confusion matrix as a heatmap
51 figure;
52 heatmap(confusionMatrix);
53 title('Confusion Matrix');
54 xlabel('Predicted Class');
55 ylabel('True Class');
56
57 % Calculate precision, recall, and F1 score for each class
58 precision = diag(confusionMatrix) ./ sum(confusionMatrix, 2);
59 recall = diag(confusionMatrix) ./ sum(confusionMatrix, 1)';
60 f1Score = 2 * (precision .* recall) ./ (precision + recall);
61
62 % Display the precision, recall, and F1 score
63 fprintf('Precision: %f\n', mean(precision));
64 fprintf('Recall: %f\n', mean(recall));
65 fprintf('F1 Score: %f\n', mean(f1Score));
66
67 % Display the precision, recall, and F1 score for each class
68 disp("precision" + precision);
69 disp("recall" + recall);
70 disp("f1Score" + f1Score);
71
72 end
73
74 % Function to define the LeNet-5 architecture
75 function [layers, options] = lenet_architecture(image_size, num_classes,
76 validation_data)
77 % Define the LeNet-5 architecture
78 layers = [
79     imageInputLayer([image_size 1], 'Name', 'input')
80
81     convolution2dLayer(5, 6, 'Padding', 'same', 'Name', 'conv_1')
82     averagePooling2dLayer(2, 'Stride', 2, 'Name', 'avgpool_1')
83
84     convolution2dLayer(5, 16, 'Padding', 'same', 'Name', 'conv_2')
85     averagePooling2dLayer(2, 'Stride', 2, 'Name', 'avgpool_2')
86
87     fullyConnectedLayer(120, 'Name', 'fc_1')
88     fullyConnectedLayer(84, 'Name', 'fc_2')
89
90     fullyConnectedLayer(num_classes, 'Name', 'fc_3')
91     softmaxLayer('Name', 'softmax')
92     classificationLayer('Name', 'output')
93 ];

```

```

94 % Specify the training options
95 options = trainingOptions('sgdm', ...
96     'InitialLearnRate',0.0001, ...
97     'MaxEpochs',10, ...
98     'Shuffle','every-epoch', ...
99     'ValidationData', validation_data, ...
100    'ValidationFrequency',30, ...
101    'Verbose',false, ...
102    'Plots','training-progress', ...
103    'ExecutionEnvironment', 'gpu');
104 end
105
106 % Function to define a custom architecture
107 function [layers, options] = custom_architecture(image_size, num_classes,
108     validation_data)
109     % Define the custom architecture
110     filterSize = 5;
111     numFilters = 16;
112     fullyConnectedLayerSize = 64;
113     l2_reg = 0.0001;
114
115     layers = [
116         imageInputLayer([image_size 1])
117
118         % Hidden Layer 1
119         convolution2dLayer(filterSize, numFilters, 'Padding', 'same')
120         batchNormalizationLayer
121         reluLayer
122         maxPooling2dLayer(2, 'Stride', 2)
123
124         % Hidden Layer 2
125         convolution2dLayer(filterSize, numFilters, 'Padding', 'same')
126         batchNormalizationLayer
127         reluLayer
128         maxPooling2dLayer(2, 'Stride', 2)
129
130         fullyConnectedLayer(fullyConnectedLayerSize, ...
131             'WeightL2Factor', l2_reg, 'BiasL2Factor', l2_reg)
132         dropoutLayer(0.2)
133
134         fullyConnectedLayer(num_classes)
135         softmaxLayer
136
137         classificationLayer
138     ];
139
140     % Specify the training options
141     options = trainingOptions('adam', ...
142         'InitialLearnRate',0.001, ...

```

```
142 'MaxEpochs',10, ...
143 'Shuffle','every-epoch', ...
144 'ValidationData', validation_data, ...
145 'ValidationFrequency',30, ...
146 'Verbose',false, ...
147 'Plots','training-progress', ...
148 'ExecutionEnvironment', 'gpu');
149
150 end
```

References

- [1] S. Basar, M. Ali, G. Ochoa-Ruiz, M. Zareei, A. Waheed, and A. Adnan, “Unsupervised color image segmentation: A case of rgb histogram based k-means clustering initialization,” *PLoS ONE*, vol. 15, no. 10, p. e0240015, 2020.
- [2] A. Kardoost, “Optimizing edge detection for image segmentation with multicut penalties,” in *Pattern Recognition(DAGM GCPR 2022)*, ser. Lecture Notes in Computer Science, vol. 13485. Springer, 2022.
- [3] C. Xiao, Q. Cao, Y. Zhong, L. Lan, X. Zhang, Z. Luo, and D. Tao, “Motiontrack: Learning motion predictor for multiple object tracking,” *arXiv preprint arXiv:2306.02585*, 2023. [Online]. Available: <https://doi.org/10.48550/arXiv.2306.02585>
- [4] T. Biswas, D. Bhattacharya, G. Mandal, and T. Das, “Motion detection using three frame differencing and cnn,” in *Computational Intelligence in Communications and Business Analytics*, ser. Communications in Computer and Information Science, K. Dasgupta, S. Mukhopadhyay, J. Mandal, and P. Dutta, Eds., vol. 1955, 2024. [Online]. Available: https://doi.org/10.1007/978-3-031-48876-4_5
- [5] D. Reynolds, *Gaussian Mixture Models*. Springer, 2015. [Online]. Available: https://doi.org/10.1007/978-1-4899-7488-4_96
- [6] M. Gröger, V. Borisov, and G. Kasneci, “Boxshrink: From bounding boxes to segmentation masks,” *arXiv preprint arXiv:2208.03142*, 2022. [Online]. Available: <https://doi.org/10.48550/arXiv.2208.03142>
- [7] Y. LeCun, L. Bottou, Y. Bengio, and P. Haffner, “Gradient-based learning applied to document recognition,” *Proceedings of the IEEE*, 1998. [Online]. Available: <https://doi.org/10.1109/5.726791>
- [8] R. A. Waelen, “The ethics of computer vision: an overview in terms of power,” *AI and Ethics*, 2023. [Online]. Available: <https://doi.org/10.1007/s43681-023-00272-x>